
SENSE: Semantic Neural Speech Synthesis from Brain Dynamics via Spatial Graph Encoding

Jisoo Park^{1*} Seonghak Lee¹ Hyojin Park² Junseok Kwon¹

¹Chung-Ang University, South Korea

²University of Birmingham, United Kingdom
susiehome@cau.ac.kr

Abstract

Reconstructing speech from non-invasive brain signals offers a promising pathway for restoring communication in individuals who are cognitively intact but unable to speak. Existing EEG-to-speech approaches formulate this task as *acoustic reconstruction*, optimizing waveform fidelity while ignoring whether the generated speech preserves high-level semantic content. In this work, we revisit this formulation and argue that EEG signals carry not only acoustic but also semantic information. We identify two key limitations of prior methods: (1) the neglect of spatial relationships between EEG electrodes, and (2) the failure to exploit the semantic structure of the N400 paradigm, where congruent and incongruent trials reflect distinct semantic processing. We propose **SENSE** (Semantic-EEG Neural Speech SynthEsis), which combines a graph-based EEG encoder over electrode geometry with EEG Semantic Conditioning (ESC), aligning EEG to a pretrained semantic space using only congruent trials. On the N400 dataset, SENSE consistently outperforms prior methods on both acoustic and semantic metrics, and model-internal channel attribution suggests distributed reliance on auditory, sensorimotor, and centro-parietal regions, consistent with known speech-perception neuroscience. In the unseen-subject setting, SENSE trained on only two subjects already surpasses the strongest baseline trained on all eighteen subjects in word error rate, and matches it on acoustic metrics with as few as eight subjects. [†]

1 Introduction

Restoring speech communication from brain activity is a central goal of brain-computer interface (BCI) research, with direct implications for individuals who are cognitively intact but unable to speak due to conditions such as amyotrophic lateral sclerosis (ALS) or locked-in syndrome [1, 2]. Beyond speech, neural decoding has enabled motor intention decoding [3, 4], visual stimulus reconstruction [5, 6], and decoding of high-level semantic representations [7, 8]. State-of-the-art results, however, largely rely on invasive or high-fidelity modalities such as electrocorticography (ECoG) and functional MRI (fMRI), which are limited by surgical requirements or low temporal resolution [9]. Among non-invasive alternatives, electroencephalography (EEG) is particularly attractive due to its portability, low cost, and millisecond-level temporal resolution [10–12].

Despite EEG’s well-known limitations of noise, artifact susceptibility, and indirect reflection of neural sources [12, 14, 15], recent work has shown that speech can be reconstructed from EEG in *perceived speech* settings, where neural activity recorded during passive listening is paired with the speech signal. Early approaches predict acoustic features and then synthesize speech from them [16], while more recent methods directly generate waveforms from EEG end-to-end [13, 17]. These models are

^{*}Work done during a research visit to the University of Birmingham.

[†]Code and audio samples are available at <https://jisoo-o.github.io/website/projects/SENSE/>




Figure 1: **Motivation for the proposed SENSE framework in EEG-to-Speech Generation.** EEG during speech perception carries acoustic, sensorimotor, and lexical-semantic information distributed across the scalp (center). Prior work [13] captures channel patterns without clear anatomical structure (left); SENSE recovers attribution aligned with all three functional regions (right) via *EEG Semantic Conditioning* and *Graph-based EEG Encoding*.

trained with waveform- or spectral-level objectives and evaluated using acoustic measures such as Mel Cepstral Distortion and Mel-Correlation [13, 17]. While these metrics capture signal fidelity, they offer limited insight into whether the reconstructed speech preserves semantic content. As a result, EEG-to-speech generation has been largely framed as an acoustic reconstruction problem, leaving open whether semantic information in EEG is being captured at all.

These observations suggest that the acoustic-only framing is incomplete. First, FESDE [17] reports that using all EEG channels outperforms restricting input to the temporal and parietal regions associated with auditory processing [18]. If the task were purely acoustic, auditory-related regions should suffice; the improvement from broader channel usage suggests that EEG-to-speech generation leverages information beyond auditory processing, though the underlying mechanism remains unclear. Second, existing models overlook the data’s structure. The standard benchmark, the N400 corpus [19], contains sentences with semantically congruent or incongruent endings designed to elicit the N400 event-related potential, a well-established index of lexical and conceptual semantic processing [20, 21]. However, prior models [13, 17] treat congruent and incongruent trials identically across all training objectives, missing the opportunity to use the congruency distinction for semantic alignment. Taken together, EEG-to-speech generation involves both acoustic and semantic information, and that the paradigm’s semantic structure could be used to ground EEG representations in linguistic meaning.

We address these gaps along two axes. First, for semantic grounding, we align EEG representations with a pretrained CLIP space [22], restricted to congruent trials so that the alignment target reflects the brain’s coherent processing of the target word. CLIP is well-suited to this setting because the N400 paradigm fixes the sentence frame and varies only the final word, concentrating semantic differences at the sentence-final position; its sentence-level embedding therefore captures variations that directly correspond to the semantic content being processed by the brain. Although CLIP is not trained on neural signals, prior work has shown that its embedding space captures structure also represented in the brain [7, 8, 23]. Second, for spatial structure, we model EEG as a graph defined by electrode geometry, in contrast to existing EEG-to-speech encoders that treat channels as unordered feature dimensions [13, 16, 17, 24]. Graph-based EEG models have been explored for classification tasks such as emotion recognition [25, 26], but to our knowledge, have not been applied to the EEG-to-speech generation. Together, these design choices yield representations that jointly capture neural activity across multiple functionally distinct regions and high-level semantic content, enabling the model to leverage information beyond auditory regions, as we confirm in Section 4.4.

We propose **SENSE** (Semantic-EEG Neural Speech SynthEsis), a framework for EEG-to-speech generation that integrates spatial modeling with paradigm-aware semantic conditioning. SENSE includes a graph-based EEG encoder that reflects electrode geometry and an EEG Semantic Conditioning (ESC) module that aligns EEG representations with a pretrained semantic space using congruent trials, incorporating semantic information into speech generation. Figure 1 illustrates the motivation of our work. We summarize our contributions as follows:

- We reformulate EEG-to-speech generation to incorporate both acoustic and semantic information, showing that semantic structure benefits both training and evaluation.
- We propose **SENSE**, a framework that combines a graph-based EEG encoder with EEG Semantic Conditioning (ESC), leveraging the congruency structure of N400-style corpora to align EEG representations with a pretrained semantic space.
- We show that SENSE improves both reconstruction quality and semantic consistency, achieving strong gains in unseen-subject generalization, and surpassing the strongest baseline trained on all 18 subjects in word error rate even when trained on only two subjects.

2 Related Work

Neural Speech and Language Decoding. Decoding speech and language from brain activity has advanced across both invasive and non-invasive modalities. Intracortical and ECoG-based approaches achieve strong performance, including continuous speech synthesis [27], high-throughput neuroprostheses [1, 2, 28], and low-resource ECoG-to-speech synthesis via wav2vec 2.0 [29] transfer learning [30], but all require invasive recordings [9]. Non-invasive alternatives are safer but operate on weaker signals: fMRI enables language reconstruction with language-model priors [7], while MEG and EEG offer high temporal resolution [31]. Several recent non-invasive works target *text* rather than waveforms (via contrastive retrieval against wav2vec 2.0 [31], discrete EEG codex with pretrained BART [32], or Whisper-based MEG-to-text decoding [33]) and thus differ from our task in either modality or output. Within EEG-to-speech *waveform generation*, early work predicted acoustic features [16], while recent end-to-end models reconstruct waveforms directly [13, 17]; these are the comparable systems on the N400 corpus [19] and serve as our baselines. All formulate the task as acoustic reconstruction, despite evidence that using all channels improves performance beyond auditory regions [17, 18]. In contrast, we treat EEG-to-speech generation as involving both acoustic and semantic information, and incorporate semantic structure into both training and evaluation.

Semantic Representations from Brain Activity. Semantic representations capture high-level meaning beyond sensory details. Prior work shows that such information is encoded in brain activity: semantic organization across cortex has been mapped during natural speech [23], and conceptual meaning can be decoded from fMRI [8]. The N400 component further indicates that semantic processing is accessible at the EEG level [20, 21]. More recently, brain signals have been aligned with pretrained embedding spaces: fMRI and EEG have been mapped into the CLIP [22] for visual retrieval and reconstruction [5, 34–37], and MEG/EEG have been aligned with wav2vec 2.0 [29] to identify corresponding speech segments from brain activity [31]. These results suggest that pretrained embedding spaces, despite not being trained on neural data, can serve as effective targets for extracting semantically meaningful information from brain signals via alignment. Unlike alignment-based approaches focused on retrieval or identification, we use semantic alignment as a conditioning mechanism within a generative pipeline that directly synthesizes speech waveforms from EEG.

EEG Representation Learning. EEG representation learning is challenging due to low signal-to-noise ratio and inter-subject variability, and limited data [12, 15, 38]. Convolutional models such as EEGNet [39] treat channels as feature dimensions processed by temporal and spatial filters and existing EEG-to-speech models follow this convention without explicitly modeling spatial relationships [13, 16, 17, 24]. Graph-based approaches model electrode geometry with learned or biologically motivated adjacencies [25, 26], but are primarily applied to classification tasks such as emotion recognition. Large-scale self-supervised pretraining has also been explored to improve generalization [40, 41]. In contrast, we apply graph-based spatial modeling to a generative setting rather than classification, and incorporate structure not through within-modality pretraining but via alignment with a semantic space learned from large-scale multimodal data.

3 Method

Figure 2 illustrates the overall pipeline. Each training example consists of an $\mathbf{X} \in \mathbb{R}^{128 \times T}$, the corresponding sentence transcript c , and the recorded speech audio y . EEG is encoded by the EEG Module (Section 3.1) into a latent representation $\mathbf{M} \in \mathbb{R}^{d \times T'}$ ($d=192$), which is then mapped to the prior of a VITS speech decoder [42], comprising a normalizing flow and a HiFi-GAN vocoder [43]. The decoder is supervised using y through the standard VITS training objectives. During training, $\mathbf{M}$




Figure 2: **SENSE training pipeline.** EEG signal $\mathbf{X}$ is encoded by a structure-aware EEG encoder into a latent $\mathbf{M}$. $\mathbf{M}$ is mapped to the VITS prior via a connector module and conditioned at the decoder using an EEG embedding. The EEG embedding is aligned with a frozen CLIP text embedding (ESC, congruent trials only). Solid arrows indicate the forward data flow, while dashed arrows represent loss supervision, connecting model predictions to their corresponding training targets.

is further guided by two supervisory signals derived from c : (1) phoneme sequence supervision, where phoneme targets are obtained by tokenizing c (Section 3.2), and (2) EEG Semantic Conditioning, where c is used as input to a frozen CLIP model (Section 3.3). At inference time, only EEG is required. A semantic conditioning vector $\mathbf{g}_{\text{eeg}} \in \mathbb{R}^d$, derived from $\mathbf{M}$, is used to directly condition the decoder.

SENSE addresses two limitations identified in Section 1. First, prior EEG-to-speech models treat EEG channels as unordered feature dimensions, ignoring their spatial structure; we address this with a graph-based encoder defined over electrode geometry. Second, prior models use the N400 corpus in a paradigm-agnostic manner, treating congruent and incongruent trials identically; we address this with EEG Semantic Conditioning (ESC), which restricts semantic alignment to congruent trials.

3.1 Structure-Aware EEG Encoder

Prior EEG-to-speech encoders are typically trained solely with a self-reconstruction or phoneme objectives [13, 17, 24], producing representations optimized primarily for acoustic fidelity. In contrast, SENSE jointly supervises the EEG module with both phonemic and semantic objectives, yielding a latent representation that captures acoustic structure, phonetic content, and high-level semantic meaning. The module consists of a spatial encoder that exploits electrode geometry, a temporal encoder that models long-range dependencies, and two auxiliary components used only during training.

Spatial Encoder. EEG electrodes are arranged according to standard scalp positions, and signals originating from spatially proximate cortical sources exhibit strong correlations [14]. Conventional 1D CNN encoders overlook this spatial structure by treating channels independently. To address this limitation, we employ a graph-based spatial encoder that explicitly models electrode topology. The encoder integrates three components: a learnable channel gate, a graph convolution over electrode geometry, and a parallel CNN skip connection.

- *Channel attention gate:* A learnable per-channel sigmoid gate, $\mathbf{X}_{\text{att}} = \mathbf{X} \odot \sigma(\mathbf{w})$, where $\mathbf{w} \in \mathbb{R}^{128}$ is a learnable parameter and σ denotes the sigmoid function, reweights each electrode prior to spatial mixing, with a positive bias on auditory-related electrodes [18] as a warm-start.

- *Graph convolution:* The 128 electrodes are represented as graph nodes with 3D coordinates defined by the BioSemi montage [44]. We construct a fixed adjacency $\mathbf{A} \in \mathbb{R}^{128 \times 128}$ using a Gaussian kernel over pairwise Euclidean distances. At each time step, node features are updated through two layers of residual graph convolution: $\mathbf{h}_i^{(\ell)} = \mathbf{h}_i^{(\ell-1)} + f^{(\ell)}\left(\sum_j A_{ij} \mathbf{h}_j^{(\ell-1)}\right)$, $\ell \in \{1, 2\}$, where each $f^{(\ell)}$ is a learnable feedforward block. The resulting node representations are then averaged over all 128 nodes and projected to dimension d , yielding $\mathbf{H}_{\text{GNN}} \in \mathbb{R}^{d \times T}$. We adopt a fixed anatomy-derived adjacency rather than learning the graph structure from data [25, 45]; with only 7,200 training pairs, joint optimization with the downstream task is prone to overfitting.

- *CNN skip connection:* Since the graph convolution involves aggregation over all 128 nodes, it can attenuate channel-specific information. To mitigate this, we introduce a parallel CNN branch applied to $\mathbf{X}_{\text{att}}$, which captures local channel interactions through learned convolutions, producing $\mathbf{H}_{\text{CNN}} \in \mathbb{R}^{d \times T}$. The outputs of two branches are combined via a learnable scalar gate $g \in \mathbb{R}$:

$\mathbf{H} = \mathbf{H}_{\text{GNN}} + \sigma(g) \cdot \tilde{\mathbf{H}}_{\text{CNN}}$, where $\tilde{\mathbf{H}}_{\text{CNN}}$ denotes $\mathbf{H}_{\text{CNN}}$ rescaled to match the variance of $\mathbf{H}_{\text{GNN}}$. The gate g is initialized to favor the graph branch during the early training.

Temporal Encoder. $\mathbf{H}$ is temporally downsampled and processed by four S4 blocks [46], yielding $\mathbf{M} \in \mathbb{R}^{d \times T'}$. We use S4 because EEG sequences remain long after downsampling; FESDE [17] adopts the same backbone for fair comparison.

Auxiliary Components. Following FESDE [17], we attach two auxiliary components used only during training. An *EEG decoder* (lightweight transposed convolutions) reconstructs the input signal from $\mathbf{M}$ under a length-masked cosine loss $\mathcal{L}_{\text{recon}}$ in (1). A *connector*, consisting of a self-attention encoder followed by a linear projection, maps $\mathbf{M}$ into the mean and log-variance of the VITS prior. A stop-gradient between the EEG module and the connector prevents VITS prior-matching gradients from dominating the EEG encoder, so $\mathbf{M}$ is shaped jointly by reconstruction, phoneme, and semantic objectives. The EEG decoder is discarded at inference; the connector remains. Implementation details are in Appendix A.

3.2 Phoneme Sequence Supervision

Mel-spectrogram reconstruction provides holistic acoustic supervision but does not explicitly encourage $\mathbf{M}$ to encode categorical phonemic structure. Prior work shows that auxiliary phoneme prediction significantly improves EEG-to-speech performance [13]. We adopt this approach with a Conformer block [47] followed by an LSTM attention decoder over a 51-symbol IPA vocabulary, trained with a CTC loss [48] against the phoneme sequence tokenized from c . The predictor is discarded at inference. Implementation details are in Appendix A.3

3.3 EEG Semantic Conditioning

Phoneme supervision encourages $\mathbf{M}$ to encode fine-grained phonetic structure, but leaves a mismatch: the decoder is conditioned on text during training, yet no such signal is available at inference. We resolve this with EEG Semantic Conditioning (ESC), which learns to map the EEG latent into the decoder’s conditioning space and replaces text-based conditioning at inference. Unlike prior brain-to-CLIP alignment work [5, 34, 36] targeting retrieval or reconstruction, ESC is designed for generative conditioning via two choices: a *congruency restriction* and a *conditioning bridge* that exposes $\mathbf{g}_{\text{eeg}}$ to the decoder during training so it remains in-distribution at inference.

Conditioning-Space Alignment. The speech decoder conditions its flow and HiFi-GAN vocoder via a learned projection $W_{\text{clip}}: \mathbf{c}_{\text{text}} \in \mathbb{R}^{512} \rightarrow \mathbf{g}_{\text{clip}} \in \mathbb{R}^d$ of CLIP text embeddings, where $\mathbf{c}_{\text{text}}$ is the output of a frozen CLIP ViT-B/32 encoder applied to the sentence c . To enable text-free conditioning, we learn an additional linear map $W_{\text{eeg}} \in \mathbb{R}^{d \times d}$ that projects the temporal mean of $\mathbf{M}$ into the same conditioning space, yielding $\mathbf{g}_{\text{eeg}} \in \mathbb{R}^d$. The alignment target is $\mathbf{g}_{\text{clip}} = \text{sg}[W_{\text{clip}}(\mathbf{c}_{\text{text}})]$, where $\text{sg}[\cdot]$ denotes stop-gradient; this prevents $\mathcal{L}_{\text{esc}}$ from pulling W_{clip} toward the EEG-reachable subspace and degrading conditioning quality, so W_{eeg} targets a fixed space shaped solely by acoustic objectives.

Congruency-Aware Cosine Loss. The N400 corpus contains congruent and incongruent sentence-final words. Aligning the EEG representation to semantic content is meaningful only for congruent trials, where neural responses reflect coherent semantic processing of the target word [20, 21]. Let $\mathcal{C}$ denote the subset of congruent indices in the batch. The ESC loss minimizes cosine distance in

the conditioning space restricted to $\mathcal{C}$: $\mathcal{L}_{\text{esc}} = \frac{1}{|\mathcal{C}|} \sum_{i \in \mathcal{C}} \left(1 - \frac{\mathbf{g}_{\text{eeg},i}^\top \mathbf{g}_{\text{clip},i}}{\|\mathbf{g}_{\text{eeg},i}\| \|\mathbf{g}_{\text{clip},i}\|} \right)$. All other training

objectives in Section 3.4 use both congruent and incongruent trials, so incongruent trials contribute to acoustic and phonemic learning without injecting noise into semantic alignment.

Conditioning Bridge. To ensure that the decoder accepts $\mathbf{g}_{\text{eeg}}$ at inference, we expose it during training. The decoder is conditioned on three sources, each sampled with roughly equal probability: $\mathbf{g}_{\text{eeg}}$, $W_{\text{clip}}(\mathbf{c}_{\text{text}})$, and a learnable null embedding (analogous to unconditional dropout in classifier-free guidance training [49], though no guided sampling is performed at inference). When conditioning on $\mathbf{g}_{\text{eeg}}$, gradients are blocked at both $\mathbf{M}$ and W_{eeg} , so acoustic gradients do not leak back into the EEG encoder and W_{eeg} is optimized solely by $\mathcal{L}_{\text{esc}}$. This ensures $\mathbf{g}_{\text{eeg}}$ is in-distribution at inference and the decoder remains compatible with text-based conditioning during training.

Table 1: **Comparison of end-to-end EEG-to-speech models** across three evaluation conditions.

	Model	MCD↓	Mel-Corr↑	STOI↑	WER↓	BERT↑	CLIP-Sim↑
Both	FESDE [17]	11.80 ± 0.12	19.87 ± 1.97	0.2785	1.2186	0.0381	0.7557
	FE-Phoneme [13]	10.52 ± 0.08	28.29 ± 1.34	0.3979	1.1610	0.0428	0.7509
	SENSE	10.44 ± 0.10	31.18 ± 0.60	0.3996	1.0186	0.0852	0.7654
Audio	FESDE [17]	11.71 ± 0.04	19.14 ± 0.47	0.2792	1.2125	0.0289	0.7601
	FE-Phoneme [13]	10.73 ± 0.06	26.53 ± 0.30	0.3918	1.1360	0.0602	0.7554
	SENSE	10.44 ± 0.01	28.50 ± 0.24	0.4039	1.0375	0.0959	0.7601
Subject	FESDE [17]	11.73 ± 0.04	18.70 ± 0.29	0.2744	1.2171	0.0269	0.7587
	FE-Phoneme [13]	10.65 ± 0.02	26.90 ± 0.08	0.3806	1.1231	0.0564	0.7555
	SENSE	9.928 ± 0.03	32.07 ± 0.14	0.4126	0.9948	0.1088	0.7647

3.4 Training Objective and Inference

Three Complementary Supervisory Signals. SENSE supervises $\mathbf{M}$ at three levels of abstraction. EEG self-reconstruction ($\mathcal{L}_{\text{recon}}$) preserves low-level signal fidelity, phoneme supervision ($\mathcal{L}_{\text{ctc}}$) imposes discrete phonetic structure, and semantic conditioning ($\mathcal{L}_{\text{esc}}$) injects high-level meaning by aligning the latent with a pretrained semantic space. These signals operate at distinct granularities (raw signal, phonemes, sentence-level meaning) and are complementary: $\mathbf{M}$ must simultaneously support EEG reconstruction, phoneme prediction, and semantic alignment.

Total Loss. We define the total loss as:

$$\mathcal{L} = \underbrace{\lambda_{\text{mel}}\mathcal{L}_{\text{mel}} + \mathcal{L}_{\text{adv}} + \mathcal{L}_{\text{fm}} + \lambda_{\text{kl}}\mathcal{L}_{\text{kl}}}_{\mathcal{L}_{\text{acoustic}} \text{ (VITS)}} + \underbrace{\lambda_{\text{recon}}\mathcal{L}_{\text{recon}}}_{\text{EEG self-recon}} + \underbrace{\lambda_{\text{ctc}}\mathcal{L}_{\text{ctc}}}_{\text{phoneme}} + \underbrace{\lambda_{\text{esc}}\mathcal{L}_{\text{esc}}}_{\text{semantic}}, \quad (1)$$

with $\{\lambda_{\text{mel}}, \lambda_{\text{kl}}, \lambda_{\text{recon}}, \lambda_{\text{ctc}}, \lambda_{\text{esc}}\} = \{45, 1.0, 1.0, 0.3, 0.5\}$. $\mathcal{L}_{\text{mel}}$ (mel-spectrogram L1 reconstruction), $\mathcal{L}_{\text{adv}}$, $\mathcal{L}_{\text{fm}}$, and $\mathcal{L}_{\text{kl}}$ follow VITS [42]. $\mathcal{L}_{\text{recon}}$ is the cosine reconstruction loss applied to the EEG decoder output (Section 3.1), $\mathcal{L}_{\text{ctc}}$ is the phoneme prediction loss of FE-Phoneme [13] (Section 3.2), and $\mathcal{L}_{\text{esc}}$ is the congruency-restricted alignment loss defined in Section 3.3.

Inference. The EEG decoder, phoneme predictor, and CLIP encoder are discarded. Raw EEG is processed by the EEG encoder to produce $\mathbf{M}$, which is mapped to $\mathbf{g}_{\text{eeg}}$ via W_{eeg} and to the VITS prior via the Connector. The decoder synthesizes the waveform conditioned on $\mathbf{g}_{\text{eeg}}$, without text or guided sampling; the null embedding from training is unused at inference and serves only as regularization.

4 Experiments

Dataset. We evaluated on the N400 EEG corpus [19], which comprises 24 subjects, each with 440 sentences, recorded using 128-channel EEG at 512 Hz. Following [13, 17], we excluded four unreliable subjects, leaving 20 subjects total, and used a standard preprocessing pipeline; speech was downsampled to 22,050 Hz with 80-band mel-spectrogram targets (Appendix A.1). The training set contains 7,200 EEG-speech pairs from 18 subjects $\times$ 400 sentences. We evaluated performance under three test conditions following [17]: *Unseen audio* (training subjects, held-out sentences; 720 pairs), *Unseen subject* (held-out subjects sub-23, sub-24 with training sentences; 800 pairs), and *Unseen both* (held-out subjects with held-out sentences; 80 pairs). For cross-subject scaling experiments, we used a nested subset structure $k=2 \subset k=4 \subset k=8 \subset k=18$ (Appendix A.5).

Baselines and Metrics. We compared SENSE against FESDE [17] and FE-Phoneme [13], the comparable existing end-to-end EEG-to-speech models on this corpus, retrained using their official implementations. Acoustic fidelity was measured by MCD, Mel-Correlation, and STOI; linguistic content by WER (via Whisper-large [50]); rescaled BERTScore (BERT) [51], and CLIP-Sim [22]. WER may exceed 1 due to insertion errors. Since incongruent trials are designed to violate the sentence-final semantic expectation, semantic metrics were computed only on congruent trials; acoustic metrics used all trials.

Training Setup. SENSE was trained for 200k iterations until validation performance plateaus, with loss weights $\lambda_{\text{ctc}} = 0.3$ and $\lambda_{\text{esc}} = 0.5$ selected via the sensitivity analysis in Appendix A.13. Baselines were trained for 100k iterations as specified in their official implementations and original papers, with extended training providing no further gains. All results were averaged over five seeds. Implementation details are in Appendix A.




Figure 3: **Qualitative comparison of generated speech.** Mel-spectrograms and Whisper-large transcriptions for ground truth, SENSE, FESDE, and FE-Phoneme. Words matching the ground truth are shown in **bold**, while semantically or phonetically similar words are underlined.

— SENSE (Ours) — FE-phoneme — FESDE — FE-phoneme @ $k=18$ (reference) ○ First k : SENSE $\geq$ FE-phoneme ($k=18$)




Figure 4: **Scaling with the number of training subjects.** Models are trained on $k \in \{2, 4, 8, 18\}$ subjects and evaluated on the held-out pair (sub-23, sub-24). Shaded regions: ± 1 std over 5 seeds. Dashed line: FE-Phoneme at $k=18$ (best baseline). Orange circles: the smallest k at which SENSE first matches or exceeds this reference.




Figure 5: **Per-subject results across various held-out pairs.** All models are retrained from scratch on the remaining 18 subjects for each pair. Error bars: ± 1 std over 5 seeds.

4.1 Main Results

SENSE outperforms both baselines across all splits and metrics (Table 1), with the largest gains on the *unseen subject* split, indicating that it learns representations that generalize across subjects rather than relying on subject-specific cues. Absolute WER values near 1 nonetheless show that intelligibility-level reconstruction from EEG remains an open challenge. BERTScore and CLIP-Sim differences are smaller, as these sentence-level metrics are less sensitive to fine-grained lexical differences than WER. Qualitative examples in Figure 3 show that SENSE produces speech that is both acoustically consistent and semantically meaningful, while baselines often fail to recover lexical content. *For additional qualitative evidence, please refer to the [audio samples](#) in the supplementary video and Appendix A.8*

4.2 Cross-Subject Generalization

We evaluated cross-subject generalization in two complementary ways: (1) scaling the number of training subjects, and (2) varying which subjects are held out. Figure 4 shows performance as the number of training subjects k increases, evaluated on a fixed held-out pair (sub-23, sub-24). Our SENSE consistently outperforms both baselines, with the largest gains in the low-data regime. Notably, SENSE already surpasses FE-Phoneme’s $k=18$ WER at $k=2$, and matches or exceeds it on MCD, Mel-Correlation, and STOI by $k=8$. Since evaluation is conducted on unseen subjects, these

Table 2: **Ablation study of the proposed SENSE components.** Each row corresponds to removing one component from the full model. w shuffled ESC uses the same ESC loss and conditioning bridge as SENSE, but randomly permutes CLIP targets within each batch.

Model	Ch.Att	GNN	CNN Skip	CTC	ESC	MCD↓	Mel-Corr↑	STOI↑	WER↓	BERT↑	CLIP-Sim↑
SENSE	✓	✓	✓	✓	✓	9.93 ±0.03	32.07 ±0.14	0.4126	0.9948	0.1088	0.7647
w/o ESC	✓	✓	✓	✓	-	10.34 ±0.03	29.43 ±0.27	0.4016	1.0323	0.0763	0.7601
w shuffled ESC	✓	✓	✓	✓	✓ (shuf)	10.38 ±0.04	29.25 ±0.28	0.4002	1.0385	0.0721	0.7595
w/o CTC	✓	✓	✓	-	✓	10.50 ±0.01	27.24 ±0.27	0.3721	1.0488	0.0588	0.7624
w/o CNN Skip	✓	✓	-	✓	✓	10.40 ±0.01	29.64 ±0.25	0.3970	1.0306	0.0724	0.7579
w/o GNN	✓	-	-	✓	✓	10.11 ±0.02	29.54 ±0.22	0.3951	1.0355	0.0788	0.7577
w/o Ch.Att	-	✓	✓	✓	✓	10.24 ±0.01	29.08 ±0.16	0.4046	1.0170	0.0926	0.7639

Table 3: **Effect of adding ESC to baseline encoders.** We incorporate ESC into FESDE and FE-Phoneme without other changes. While both baselines improve, they remain consistently below SENSE across all metrics, indicating that ESC complements rather than replaces the SENSE encoder.

Backbone	ESC	MCD ↓	Mel-Corr ↑	STOI ↑	WER ↓	BERT ↑	CLIP-sim ↑
FESDE [17]	✓	10.93 ±0.03	26.27 ±0.19	0.363	1.084	0.045	0.7500
FE-Phoneme [13]	✓	10.04 ±0.02	31.17 ±0.21	0.411	1.102	0.070	0.7544
SENSE	✓	9.93 ±0.03	32.07 ±0.14	0.413	0.995	0.109	0.7647

Table 4: **Performance across different graph adjacency structures.** *BioSemi topo*: anatomy-based adjacency derived from electrode geometry (used in SENSE). *Fully connected*: uniform adjacency. *Random*: fixed random adjacency. *GAT*: learned edge attention via Graph Attention Network.

Model	Graph Topology	CNN Skip	MCD↓	Mel-Corr↑	STOI↑	WER↓
SENSE	Biosemi topo	✓	9.93 ±0.03	32.07 ±0.14	0.4126	0.9948
GNN (topo, no skip)	Biosemi topo	-	10.40 ±0.02	29.64 ±0.25	0.3970	1.0306
GNN (GAT attn. topo)	Biosemi topo	-	10.21 ±0.03	29.10 ±0.10	0.3835	1.0183
GNN (fully-conn.)	Fully-conn.	-	10.96 ±0.05	26.17 ±0.27	0.3684	1.0945
GNN (random)	Random	-	10.67 ±0.03	27.15 ±0.18	0.3803	1.0806

Table 5: **Performance across different semantic encoders in the ESC module.**

Semantic Encoder	MCD ↓	Mel-Corr ↑	STOI ↑	WER ↓	BERT ↑
SBERT [52]	10.06 ±0.01	30.25 ±0.09	0.399	1.011	0.094
CLAP [53]	10.02 ±0.04	29.62 ±0.21	0.388	1.013	0.092
CLIP [22]	9.93 ±0.03	32.07 ±0.14	0.413	0.995	0.109

results reflect cross-subject generalization rather than in-distribution scaling, suggesting that SENSE learns representations driven by stimulus content rather than subject-specific patterns.

To test whether SENSE’s advantage depends on the choice of held-out subjects, we retrained all models from scratch using two additional held-out pairs, (sub-1, sub-2) and (sub-7, sub-8). Figure 5 reports per-subject performance across the resulting six held-out subjects. SENSE outperforms both baselines on every metric and every subject, with consistent relative gains over FE-Phoneme. Although absolute performance varies modestly across subjects, the relative ranking of the methods remains consistent. Full per-subject results are provided in Appendix A.14.

4.3 Ablation Study

Table 2 shows that all components of the proposed SENSE contribute to overall performance. Removing ESC degrades semantic metrics, while removing CTC causes an even larger increase in WER, indicating that phoneme-level supervision also preserves intelligibility (Table 2). Spatial components (GNN, CNN skip, channel attention) provide consistent gains, complementing both acoustic and semantic modeling. Adding ESC alone to baseline encoders improves them but still falls short of SENSE (Table 3), showing that semantic conditioning is insufficient without our encoder design. To verify that ESC exploits genuine EEG-semantic correspondence rather than acting as generic regularization, we trained SENSE with CLIP targets randomly shuffled within each batch (Table 2, *w shuffled ESC* row). Shuffled targets yield performance comparable to (and slightly below) the w/o ESC baseline, confirming that ESC’s gain depends on correct semantic targets rather than on auxiliary supervision alone.

We compared four graph adjacency schemes (Table 4). The anatomy-based BioSemi topology outperforms learnable GAT, fully-connected, and random adjacencies, indicating that local distance-based connectivity is the relevant inductive bias. Jointly learning the graph from only 7,200 pairs is under-constrained, while removing structure entirely is more harmful. Adding the CNN skip on top of the BioSemi GNN further improves performance, confirming complementarity between graph-based spatial mixing and channel-wise feature aggregation.




Figure 6: **EEG-text alignment across semantic encoders.** Cosine similarity matrices between mean EEG embeddings (rows) and text embeddings (columns) for matched words. Stronger word-specific alignment appears as a brighter diagonal. Δ is the gap between diagonal and off-diagonal means.




Figure 7: **ESC attribution across encoder variants.** (a) Anatomical graph used in SENSE. (b–d) Per-channel gradient magnitude of EEG-CLIP alignment, normalized per model. (b) Without GNN: centro-parietal concentration consistent with N400. (c) GNN without skip: bias toward central hub electrodes, a graph-topology artifact. (d) SENSE: tripartite pattern over temporal-auditory (TP7/TP8), motor (C3/C4), and centro-parietal (CPz). Circles: top-8 channels per model; colors: functional ROI.

We also evaluated three pretrained semantic encoders for the ESC alignment target (Table 5): CLIP, CLAP [53], and SBERT [52]. CLIP outperforms CLAP and SBERT on all metrics. CLAP, despite being trained for audio-language alignment, performs worse than text-only SBERT, possibly because the N400 stimuli often involve concrete, image-able nouns (*chair, dog, ants*) for which CLIP’s image-text pretraining provides more discriminative semantic structure. Figure 6 corroborates this: CLIP shows the brightest diagonal and the largest gap between diagonal and off-diagonal values.

4.4 Channel Attribution Analysis

We complemented the quantitative results with two analyses of the model’s internal behavior, examining whether the learned representations align with our motivation that EEG during speech perception encodes information beyond auditory regions. The channel attention gate (Section 3.1) is initialized with a positive bias toward auditory-related electrodes and zero elsewhere; no prior is placed on sensorimotor or centro-parietal regions, so the patterns below emerge from training rather than design. The *w/o Ch.Att* row in Table 2 initializes the gate without the auditory warm-start (all weights set to zero); the moderate performance drop indicates that the gate’s emergent non-auditory attribution is not merely an artifact of initialization. After training, auditory-prior channels are modestly down-weighted while central electrodes (C3, C4) and frontal-central regions gain importance, with occipital channels suppressed (Appendix A.11).

To further investigate which channels contribute to EEG-text alignment, we computed per-channel gradient magnitudes of the ESC alignment (Figure 7). We treat these gradient maps as model-internal diagnostics rather than evidence of neural sources, given known limitations of saliency-based attribution. The maps reveal three patterns: without a GNN, attribution concentrates over centro-parietal regions consistent with the N400 response. A GNN without a skip connection shifts emphasis to central hub electrodes, a topology artifact of dense connectivity rather than a physiological signal. The full SENSE model recovers a distributed pattern over temporal-auditory (TP7, TP8), central motor (C3, C4), and centro-parietal (CPz) regions. The prominence of C3/C4, which are not explicitly favored by design, is consistent with prior findings that speech perception engages sensorimotor cortex during passive listening [54–56]; we read this as the model capturing distributed information aligned with known properties of speech processing, not as direct evidence of neural generators. This supports our central hypothesis that EEG-to-speech generation benefits from non-auditory channels. Training dynamics are detailed in Appendix A.12.

5 Conclusions

We argued that EEG-to-speech should be approached as a joint acoustic-semantic problem, not as acoustic reconstruction alone. We introduced **SENSE**, which combines a graph-based spatial encoder over electrode geometry with EEG Semantic Conditioning (ESC). On the N400 dataset, **SENSE** consistently outperforms prior methods on both acoustic and semantic metrics, with particularly strong cross-subject generalization. Channel attribution further recovers patterns spanning auditory, sensorimotor, and centro-parietal regions consistent with neuroscience of speech-perception. However, reliable word-level reconstruction remains challenging, as indicated by WER values near 1, and full waveform synthesis has only been evaluated on the controlled N400 corpus despite promising representation-level adaptation to naturalistic speech. Channel attribution remains model-internal rather than neural-source localization, and evaluation on larger naturalistic and active-speech datasets is an important next step.

5.1 Broader Impact

This work uses a previously released public EEG benchmark collected under prior ethical approval, with no new participant recruitment, neural data collection, or sensing hardware, and should not be interpreted as unrestricted thought decoding. Active listening and laboratory conditions are not reliable safeguards, and future applications should protect neural data through measures such as on-device processing, explicit opt-in consent, restrictions on unauthorized access, and deployment-specific oversight and governance. These considerations follow the UNESCO Recommendation on the Ethics of Neurotechnology (2025), including its guidance on protecting neural data and information that may enable inferences about mental states.

Acknowledgments

This work was supported in part by the National Research Foundation of Korea(NRF) grant funded by the Korea government(MSIT) (RS-2025-02217071), in part by the Institute of Information & communications Technology Planning & Evaluation (IITP) grant funded by the Korea government(MSIT) (RS-2021-II211341), and in part by the AI Seoul Tech Research Support Program of the Seoul Future Foundation.

References

- [1] Francis R Willett, Erin M Kunz, Chaofei Fan, Donald T Avansino, Guy H Wilson, Eun Young Choi, Foram Kamdar, Matthew F Glasser, Leigh R Hochberg, Shaul Druckmann, et al. A high-performance speech neuroprosthesis. *Nature*, 620(7976):1031–1036, 2023.
- [2] Sean L Metzger, Kaylo T Littlejohn, Alexander B Silva, David A Moses, Margaret P Seaton, Ran Wang, Maximilian E Dougherty, Jessie R Liu, Peter Wu, Michael A Berger, et al. A high-performance neuroprosthesis for speech decoding and avatar control. *Nature*, 620(7976): 1037–1046, 2023.
- [3] Dennis J McFarland and Jonathan R Wolpaw. Brain-computer interfaces for communication and control. *Communications of the ACM*, 54(5):60–66, 2011.
- [4] JD Simeral, Sung-Phil Kim, MJ Black, JP Donoghue, and LR Hochberg. Neural control of cursor trajectory and click by a human with tetraplegia 1000 days after implant of an intracortical microelectrode array. *Journal of neural engineering*, 8(2):025027, 2011.
- [5] Yu Takagi and Shinji Nishimoto. High-resolution image reconstruction with latent diffusion models from human brain activity. In *CVPR*, pages 14453–14463, 2023.
- [6] Yu-Ting Lan, Kan Ren, Yansen Wang, Wei-Long Zheng, Dongsheng Li, Bao-Liang Lu, and Lili Qiu. Seeing through the brain: image reconstruction of visual perception from human brain signals. *arXiv preprint arXiv:2308.02510*, 2023.
- [7] Jerry Tang, Amanda LeBel, Shailee Jain, and Alexander G Huth. Semantic reconstruction of continuous language from non-invasive brain recordings. *Nature Neuroscience*, 26(5):858–866, 2023.
- [8] Francisco Pereira, Bin Lou, Brianna Pritchett, Samuel Ritter, Samuel J Gershman, Nancy Kanwisher, Matthew Botvinick, and Evelina Fedorenko. Toward a universal decoder of linguistic meaning from brain activation. *Nature communications*, 9(1):963, 2018.
- [9] Markus Ullsperger and Stefan Debener. *Simultaneous EEG and fMRI: recording, analysis, and application*. 2010.
- [10] Steven J. Luck. *An Introduction to the Event-Related Potential Technique*. MIT Press, Cambridge, MA, 2nd edition, 2014.
- [11] Reza Abiri, Soheil Borhani, Eric W Sellers, Yang Jiang, and Xiaopeng Zhao. A comprehensive review of eeg-based brain–computer interface paradigms. *Journal of neural engineering*, 16(1): 011001, 2019.
- [12] Yannick Roy, Hubert Banville, Isabela Albuquerque, Alexandre Gramfort, Tiago H Falk, and Jocelyn Faubert. Deep learning-based electroencephalography analysis: a systematic review. *Journal of neural engineering*, 16(5):051001, 2019.
- [13] Jihwan Lee, Tiantian Feng, Aditya Kommineni, Sudarsana Reddy Kadiri, and Shrikanth Narayanan. Enhancing listened speech decoding from eeg via parallel phoneme sequence prediction. In *ICASSP*, 2025.
- [14] Paul L Nunez and Ramesh Srinivasan. *Electric fields of the brain: the neurophysics of EEG*. Oxford university press, 2006.
- [15] Xun Chen, Chang Li, Aiping Liu, Martin J McKeown, Ruobing Qian, and Z Jane Wang. Toward open-world electroencephalogram decoding via deep learning: A comprehensive survey. *IEEE Signal Processing Magazine*, 39(2):117–134, 2022.
- [16] Gautam Krishna, Co Tran, Yan Han, Mason Carnahan, and Ahmed H Tewfik. Speech synthesis using eeg. In *ICASSP*, pages 1235–1238, 2020.
- [17] Jihwan Lee, Aditya Kommineni, Tiantian Feng, Kleanthis Avramidis, Xuan Shi, Sudarsana Reddy Kadiri, and Shrikanth Narayanan. Toward fully-end-to-end listened speech decoding from EEG signals. In *Interspeech*, pages 2110–2114, 2024.

- [18] Rini A Sharon, Shrikanth Narayanan, Mriganka Sur, and Hema A Murthy. An empirical study of speech processing in the brain by analyzing the temporal syllable structure in speech-input induced eeg. In *ICASSP*, pages 4090–4094, 2019.
- [19] Kathryn K Toffolo, Edward G Freedman, and John J Foxe. Evoking the n400 event-related potential (erp) component using a publicly available novel set of sentences with semantically incongruent or congruent eggplants (endings). *Neuroscience*, 501:143–158, 2022.
- [20] Marta Kutas and Steven A Hillyard. Reading senseless sentences: Brain potentials reflect semantic incongruity. *Science*, 207(4427):203–205, 1980.
- [21] Marta Kutas and Kara D Federmeier. Thirty years and counting: finding meaning in the n400 component of the event-related brain potential (erp). *Annual review of psychology*, 62(1): 621–647, 2011.
- [22] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, pages 8748–8763, 2021.
- [23] Alexander G Huth, Wendy A De Heer, Thomas L Griffiths, Frédéric E Theunissen, and Jack L Gallant. Natural speech reveals the semantic maps that tile human cerebral cortex. *Nature*, 532(7600):453–458, 2016.
- [24] Aditya Kommineni, Kleanthis Avramidis, Richard Leahy, and Shrikanth Narayanan. Knowledge-guided eeg representation learning. In *EMBC*, pages 1–6, 2024.
- [25] Tengfei Song, Wenming Zheng, Peng Song, and Zhen Cui. Eeg emotion recognition using dynamical graph convolutional neural networks. *IEEE Transactions on Affective Computing*, 11(3):532–541, 2018.
- [26] Peixiang Zhong, Di Wang, and Chunyan Miao. Eeg-based emotion recognition using regularized graph neural networks. *IEEE Transactions on Affective Computing*, 13(3):1290–1301, 2020.
- [27] Gopala K Anumanchipalli, Josh Chartier, and Edward F Chang. Speech synthesis from neural decoding of spoken sentences. *Nature*, 568(7753):493–498, 2019.
- [28] David A Moses, Sean L Metzger, Jessie R Liu, Gopala K Anumanchipalli, Joseph G Makin, Pengfei F Sun, Josh Chartier, Maximilian E Dougherty, Patricia M Liu, Gary M Abrams, et al. Neuroprosthesis for decoding speech in a paralyzed person with anarthria. *New England Journal of Medicine*, 385(3):217–227, 2021.
- [29] Alexei Baevski, Yuhao Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. *NeurIPS*, 33:12449–12460, 2020.
- [30] Miseul Kim, Zhenyu Piao, Jihyun Lee, and Hong-Goo Kang. Braintalker: Low-resource brain-to-speech synthesis with transfer learning using wav2vec 2.0. In *2023 IEEE EMBS International Conference on Biomedical and Health Informatics (BHI)*, pages 1–5. IEEE, 2023.
- [31] Alexandre Défossez, Charlotte Caucheteux, Jérémy Rapin, Ori Kabeli, and Jean-Rémi King. Decoding speech perception from non-invasive brain recordings. *Nature Machine Intelligence*, 5(10):1097–1107, 2023.
- [32] Yiqun Duan, Jinzhao Zhou, Zhen Wang, Yu-Kai Wang, and Chin-teng Lin. Dewave: Discrete encoding of eeg waves for eeg to text translation. *NeurIPS*, 36:9907–9918, 2023.
- [33] Yiqian Yang, Yiqun Duan, Qiang Zhang, Hyejeong Jo, Jinni Zhou, Won Hee Lee, Renjing Xu, and Hui Xiong. Neuspeech: Decode neural signal as speech. In *ICASSP*, pages 6636–6640. IEEE, 2026.
- [34] Paul Scotti, Atmadeep Banerjee, Jimmie Goode, Stepan Shabalin, Alex Nguyen, Aidan Dempster, Nathalie Verlinde, Elad Yundler, David Weisberg, Kenneth Norman, et al. Reconstructing the mind’s eye: fmri-to-image with contrastive learning and diffusion priors. *NeurIPS*, 36: 24705–24728, 2023.

- [35] Paul S Scotti, Mihir Tripathy, Cesar Kadir Torrico Villanueva, Reese Kneeland, Tong Chen, Ashutosh Narang, Charan Santhirasegaran, Jonathan Xu, Thomas Naselaris, Kenneth A Norman, et al. Mindeye2: Shared-subject models enable fmri-to-image with 1 hour of data. *arXiv preprint arXiv:2403.11207*, 2024.
- [36] Yonghao Song, Bingchuan Liu, Xiang Li, Nanlin Shi, Yijun Wang, and Xiaorong Gao. Decoding natural images from eeg for object recognition. *arXiv preprint arXiv:2308.13234*, 2023.
- [37] Dongyang Li, Chen Wei, Shiyang Li, Jiachen Zou, Haoyang Qin, and Quanying Liu. Visual decoding and reconstruction via eeg embeddings with guided diffusion. *arXiv preprint arXiv:2403.07721*, 2024.
- [38] Fabien Lotte, Laurent Bougrain, Andrzej Cichocki, Maureen Clerc, Marco Congedo, Alain Rakotomamonjy, and Florian Yger. A review of classification algorithms for eeg-based brain-computer interfaces: a 10 year update. *Journal of neural engineering*, 15(3):031005, 2018.
- [39] Vernon J Lawhern, Amelia J Solon, Nicholas R Waytowich, Stephen M Gordon, Chou P Hung, and Brent J Lance. Eegnet: a compact convolutional neural network for eeg-based brain-computer interfaces. *Journal of neural engineering*, 15(5):056013, 2018.
- [40] Demetres Kostas, Stephane Aroca-Ouellette, and Frank Rudzicz. Bendr: Using transformers and a contrastive self-supervised learning task to learn from massive amounts of eeg data. *Frontiers in Human Neuroscience*, 15:653659, 2021.
- [41] Wei-Bang Jiang, Li-Ming Zhao, and Bao-Liang Lu. Large brain model for learning generic representations with tremendous eeg data in bci. *arXiv preprint arXiv:2405.18765*, 2024.
- [42] Jaehyeon Kim, Jungil Kong, and Juhee Son. Conditional variational autoencoder with adversarial learning for end-to-end text-to-speech. In *ICML*, pages 5530–5540. PMLR, 2021.
- [43] Jungil Kong, Jaehyeon Kim, and Jaekyoung Bae. Hifi-gan: Generative adversarial networks for efficient and high fidelity speech synthesis. *NeurIPS*, 33:17022–17033, 2020.
- [44] Alexandre Gramfort, Martin Luessi, Eric Larson, Denis A Engemann, Daniel Strohmeier, Christian Brodbeck, Roman Goj, Mainak Jas, Teon Brooks, Lauri Parkkonen, et al. Meg and eeg data analysis with mne-python. *Frontiers in Neuroinformatics*, 7:267, 2013.
- [45] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. Graph attention networks. *arXiv preprint arXiv:1710.10903*, 2017.
- [46] Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. *arXiv preprint arXiv:2111.00396*, 2021.
- [47] Anmol Gulati, James Qin, Chung-Cheng Chiu, Niki Parmar, Yu Zhang, Jiahui Yu, Wei Han, Shibo Wang, Zhengdong Zhang, Yonghui Wu, et al. Conformer: Convolution-augmented transformer for speech recognition. *arXiv preprint arXiv:2005.08100*, 2020.
- [48] Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber. Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks. In *ICML*, pages 369–376, 2006.
- [49] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- [50] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *ICML*, pages 28492–28518, 2023.
- [51] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. Bertscore: Evaluating text generation with bert. *arXiv preprint arXiv:1904.09675*, 2019.
- [52] Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *EMNLP-IJCNLP*, pages 3982–3992, 2019.

- [53] Yusong Wu, Ke Chen, Tianyu Zhang, Yuchen Hui, Taylor Berg-Kirkpatrick, and Shlomo Dubnov. Large-scale contrastive language-audio pretraining with feature fusion and keyword-to-caption augmentation. In *ICASSP*, pages 1–5, 2023.
- [54] Stephen M Wilson, Ayşe Pinar Saygin, Martin I Sereno, and Marco Iacoboni. Listening to speech activates motor areas involved in speech production. *Nature neuroscience*, 7(7):701–702, 2004.
- [55] M Florencia Assaneo and David Poeppel. The coupling between auditory and motor cortices is rate-restricted: Evidence for an intrinsic speech-motor rhythm. *Science advances*, 4(2): eaao3842, 2018.
- [56] Hyojin Park, Robin AA Ince, Philippe G Schyns, Gregor Thut, and Joachim Gross. Frontal top-down signals increase coupling of auditory low-frequency oscillations to continuous speech in human listeners. *Current Biology*, 25(12):1649–1653, 2015.
- [57] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.
- [58] Cees H Taal, Richard C Hendriks, Richard Heusdens, and Jesper Jensen. A short-time objective intelligibility measure for time-frequency weighted noisy speech. In *ICASSP*, pages 4214–4217, 2010.
- [59] Michael P Broderick, Andrew J Anderson, Giovanni M Di Liberto, Michael J Crosse, and Edmund C Lalor. Electrophysiological correlates of semantic dissimilarity reflect the comprehension of natural, narrative speech. *Current Biology*, 28(5):803–809, 2018.